

Islamic University of Technology



Report on Lab 06

(CSE 4618 Artificial Intelligence Lab)

Name: **Md Mahmudul Islam Mahin**
Student ID: **210042140**

Department: **CSE**
Program: **BSc in SWE**
Date: **8th September 2025**

Contents

1	Introduction	2
2	Problem Analysis	2
3	Solution Explanation	2
4	Findings and Insights	3
5	Challenges Faced	3
6	Hyperparameter Exploration	3
7	GitHub Repository	3

1 Introduction

This report details the implementation of a Q-learning agent for the reinforcement learning tasks outlined in Lab 6. The objective was to develop a Q-learning algorithm to learn optimal policies through interaction with environments like Gridworld, Crawler, and Pacman by implementing key methods in `qlearningAgents.py`, including `computeValueFromQValues`, `computeActionFromQValues`, `getAction`, and `update`. This report discusses the problem, solution, findings, challenges, and hyperparameter exploration.

2 Problem Analysis

The task required implementing a Q-learning agent that learns an optimal policy without prior knowledge of the environment's Markov Decision Process (MDP). Unlike value iteration, Q-learning is a model-free approach that updates Q-values based on state-action-reward transitions. The challenge was to handle exploration versus exploitation, ties in action selection, and unseen state-action pairs with a Q-value of 0, which could be optimal if known actions had negative Q-values. The implementation needed to support diverse environments (Gridworld, Crawler, Pacman) with varying state and action spaces.

3 Solution Explanation

The Q-learning agent was implemented in `qlearningAgents.py` with the following approach:

Initialization:

Used `util.Counter` to store Q-values, initialized to 0 for unseen state-action pairs.

getQValue:

Retrieves the Q-value for a state-action pair, returning 0.0 for unseen pairs.

computeValueFromQValues:

Computes the maximum Q-value over legal actions, returning 0.0 for terminal states.

computeActionFromQValues:

Selects the action with the highest Q-value, breaking ties randomly using `random.choice`.

getAction:

Implements epsilon-greedy action selection using `util.flipCoin`.

update:

Updates Q-values using the formula:

$$Q(s, a) \leftarrow (1 - \alpha) \cdot Q(s, a) + \alpha \cdot (r + \gamma \cdot \max_{a'} Q(s', a'))$$

Key decisions included using `util.Counter` for efficient storage and random tie-breaking to avoid biases. The complexity is $O(|A|)$ per update, where $|A|$ is the number of legal actions. Here is the update method:

```
1 def update(self, state, action, nextState, reward):
2     oldQ = self.getQValue(state, action)
3     nextValue = self.computeValueFromQValues(nextState)
4     self.qValues[(state, action)] = (1 - self.alpha) * oldQ + self.alpha * (reward
        + self.discount * nextValue)
```

4 Findings and Insights

In Gridworld, the agent learned the optimal policy over multiple episodes, with Q-values converging when guided along the optimal path for four episodes with noise disabled. The epsilon-greedy strategy led to early exploration of suboptimal actions, which improved long-term learning by discovering better paths, emphasizing the exploration-exploitation trade-off.

5 Challenges Faced

Ensuring random tie-breaking in `computeActionFromQValues` was challenging; initial deterministic selection caused biases, resolved by using `random.choice`. Debugging Q-value updates to match autograder expectations required testing with `noise 0.0`.

6 Hyperparameter Exploration

Performance varied with hyperparameters:

- **alpha (learning rate):** High alpha (e.g., 0.5) caused unstable Q-values, while `alpha=0.2` ensured stable but slower convergence.
- **gamma (discount factor):** `gamma=0.8` balanced immediate and future rewards, while `gamma=0.9` favored long-term rewards, risking overfitting.

7 GitHub Repository

I will be uploading the lab tasks in the following repository: [CSE 4618: Artificial Intelligence](#)